Overlapping with language modelling

Code to reproduce the results from "Alleviating Sequence Information Loss with Data Overlapping and PrimeBatch Sizes".

Hold experiments on 4 models using the overlapping:

- **awd**, AWD ASGD Weight-Dropped LSTM, (`/awd`)
- **text simple**, a very basic lstm for language modelling, (`/simple`)
- **mos**, MOS Mixture of Softmaxes, (`/mos`)
- **voice simple**, a very basic LSTM for emotion detection on voice, (`/emotions`)

To specify which model to run, use `--main-model {simple-lstm | awd-lstm | mos-lstm | emotions-simple-lstm}` argument. There are additional common paramaters, as well as specific parameters for each model. Those can be found in `main_run.py`.

The taxonomy in the code may differe a bit from the paper, especially regarding the type of experiments. Here is the corresponding terms:

In the code	In the paper
No order	Extreme TOI
Local order	Inter-batch TOI
Standard order	Standard TOI
Total order (P)	Alleviated TOI (P)

Set-up

Download the data (PTB, WT2, WT103):

```
chmod +x get_data.sh
./get_data.sh
```

For emotions, add in `data/IEMOCAP/` the `all_features_cv` files.

We use python `3.6` with Pytorch `0.4.1`. To create a new python environement and install dependencies, run:

```
python3.6 -m virtualenv venv
source venv/bin/activate
pip3 install -r requirements.txt
```

About the files

`main_run.py` is the main entry point that parses arguments, does the global initialization and runs the corresponding model and task.

`awd/`, `emotions/`, `mos/` and `simple/` are the different models directories. `common/` holds the common initilization and utilities, such as the different data iterators, which are in the `DataSelector` class in `common/excavator.py`.

The `main_run.py` file, after performing the common initilizations, imports the `main.py` file corresponding to the choosen model.

Commands to reproduce the experiments

Note: Those results do not use prime batch size, but the default parameters. To have better results, adapt the `--batch-size` param to the closest prime number.

Quick anchors navigation:

Model	Dataset	Experiments
AWD	PTB	Extreme / Inter-batch / Original / Alleviated TOI
	WT2	Extreme / Inter-batch / Original / Alleviated TOI
	WT103	Extreme / Inter-batch / Original / Alleviated TOI
Text simple LSTM	PTB	Extreme / Inter-batch / Original / Alleviated TOI
	WT2	Extreme / Inter-batch / Original / Alleviated TOI
MOS	PTB	Original / Alleviated TOI
Voice simple LSTM	IEMOCAP	Extreme / Inter-batch / Original / Alleviated TOI

AWD PTB

Extreme TOI:

Expected results: `66.38` / `63.49` (validation / testing)

```
python3 main_run.py --main-model awd-lstm --batch-size 20 --data data/penn
--dropouti 0.4 --dropouth 0.25 --seed 141 --seed-shuffle 141 --epochs 1000
--shuffle-full-seq
```

Inter-batch TOI:

Expected results: `66.96` / `64.20` (validation / testing)

```
python3 main_run.py --main-model awd-lstm --batch-size 20 --data data/penn
--dropouti 0.4 --dropouth 0.25 --seed 141 --seed-shuffle 141 --epochs 1000
--shuffle-row-seq
```

Standard TOI:

Expected results: 61.28 / 58.94 (validation / testing)

```
python3 main_run.py --main-model awd-lstm --batch-size 20 --data data/penn
--dropouti 0.4 --dropouth 0.25 --seed 141 --epochs 1000
```

Alleviated TOI {2,5,7,10}:

Expected results (validation / testing):

- 2: 61.73 / 59.37
- 5: 63.37 / 60.50
- 7: 59.22 / 56.7
- 10: 68.09 / 65.88

```
overlaps=(2 5 7 10)
epochs=1000
for k in "${overlaps[@]}"
do
    :
    python3 main_run.py --main-model awd-lstm --batch-size 20 --data
data/penn --dropouti 0.4 --dropouth 0.25 --seed 141 --epochs
"$(($epochs/$k))" --init-seq "overlapCN_${k}"
    sleep 10
done
```

AWD WT2

Extreme TOI

Expected results: 77.14 / 73.52 (validation / testing)

```
python3 main_run.py --main-model awd-lstm --epochs 750 --data
/data/noemien.kocher/datasets/wikitext-2 --dropouth 0.2 --seed 1882 --
batch-size 80 --shuffle-full-seq
```

Inter-batch TOI

Expected results: 76.08 / 72.61 (validation / testing)

```
python main_run.py --main-model awd-lstm --epochs 750 --data
/data/noemien.kocher/datasets/wikitext-2 --dropouth 0.2 --seed 1882 --
batch-size 80 --shuffle-row-seq
```

Standard TOI

Expected results: 68.50 / 65.86 (validation / testing)

```
python3 main_run.py --main-model awd-lstm --epochs 750 --data
/data/noemien.kocher/datasets/wikitext-2 --dropouth 0.2 --seed 1882 --
batch-size 80
```

Alleviated TOI {2,5,7,10}

Expected results (validation / testing):

- 2: 68.56 / 65.51
- 5: 69.56 / 66.33
- 7: 67.48 / 64.87
- 10: 72.95 / 69.69

```
overlaps=(2 5 7 10)
epochs=750
for k in "${overlaps[@]}"
do
    :
    python3 main_run.py --main-model awd-lstm --data
/data/noemien.kocher/datasets/wikitext-2 --dropouth 0.2 --seed 1882 --
batch-size 80 --epochs "$(($epochs/$k))" --init-seq "overlapCN_${k}"
    sleep 10
done
```

AWD WT103

Extreme TOI

Expected results: 35.22 / 36.19 (validation / testing)

```
python3 -u main_run.py --main-model awd-lstm --epochs 14 --nlayers 4 --
emsize 400 --nhid 2500 --alpha 0 --beta 0 --dropoute 0 --dropouth 0.1 --
dropouti 0.1 --dropout 0.1 --wdrop 0 --wdecay 0 --bptt 140 --batch-size 60
--optimizer adam --lr 1e-3 --data /data/noemien.kocher/datasets/wikitext-
103 --when 12 --model QRNN --shuffle-full-seq
```

Inter-batch TOI

Expected results: 35.41 / 36.39 (validation / testing)

```
python3 -u main_run.py --main-model awd-lstm --epochs 14 --nlayers 4 --
emsize 400 --nhid 2500 --alpha 0 --beta 0 --dropoute 0 --dropouth 0.1 --
dropouti 0.1 --dropout 0.1 --wdrop 0 --wdecay 0 --bptt 140 --batch-size 60
--optimizer adam --lr 1e-3 --data /data/noemien.kocher/datasets/wikitext-
103 --when 12 --model QRNN --shuffle-row-seq
```

Standard TOI

Expected results: 32.18 / 32.94 (validation / testing)

```
python3 -u main_run.py --main-model awd-lstm --epochs 14 --nlayers 4 --
emsize 400 --nhid 2500 --alpha 0 --beta 0 --dropoute 0 --dropouth 0.1 --
dropouti 0.1 --dropout 0.1 --wdrop 0 --wdecay 0 --bptt 140 --batch-size 60
--optimizer adam --lr 1e-3 --data /data/noemien.kocher/datasets/wikitext-
103 --when 12 --model QRNN
```

Alleviated TOI {2,5,7,10}

Expected results (validation / testing):

- 2: 36.94 / 34.31
- 5: 38.50 / 40.04
- 7: 31.78 / 32.72
- 10: 48.28 / 49.49

```
# base num epochs is 14
overlaps=(2 5 7 10)
when_steps=147456
max_steps=172032
for i in "${!overlaps[@]}"
do
        :
        python3 -u main_run.py --main-model awd-lstm --epochs 14 --nlayers
4 --emsize 400 --nhid 2500 --alpha 0 --beta 0 --dropoute 0 --dropouth 0.1
--dropouti 0.1 --dropout 0.1 --wdrop 0 --wdecay 0 --bptt 140 --batch-size
60 --optimizer adam --lr 1e-3 --data
/data/noemien.kocher/datasets/wikitext-103 --when-steps "$when_steps" --
model QRNN --init-seq "overlapCN_${overlaps[$i]}" --log-dir
/data/noemien.kocher/logs/ --max-steps "$max_steps"
        sleep 10
done
```

Simple PTB

Extreme TOI:

Expected results: 81.97 / 79.08 (validation / testing)

```
python3 main_run.py --main-model simple-lstm --epochs 100 --batch-size 20
--dropout 0.15 --nlayers 2 --bptt 70 --nhid 1500 --lr-decay 1 --shuffle-
full-seq
```

Inter-batch TOI:

Expected results: 81.67 / 78.59 (validation / testing)

```
python3 main_run.py --main-model simple-lstm --epochs 100 --batch-size 20
--dropout 0.15 --nlayers 2 --bptt 70 --nhid 1500 --lr-decay 1 --shuffle-
row-seq
```

Standard TOI:

Expected results: 77.54 / 75.36 (validation / testing)

```
python3 main_run.py --main-model simple-lstm --epochs 100 --batch-size 20
--dropout 0.15 --nlayers 2 --bptt 70 --nhid 1500 --lr-decay 1
```

Alleviated TOI {2,5,7,10}:

Expected results (validation / testing):

- 2: 78.48 / 76.55
- 5: 91.95 / 89.64
- 7: 77.47 / 74.98
- 10: 92.92 / 92.07

```
overlaps=(2 5 7 10)
epochs=100
for k in "${overlaps[@]}"
do
    :
    python3 main_run.py --main-model simple-lstm --epochs
"$(($epochs/$k))" --batch-size 20 --dropout 0.15 --nlayers 2 --bptt 70 --
nhid 1500 --lr-decay 1 --init-seq "overlapCN_${k}"
    sleep 10
done
```

Simple WT2

Extreme TOI

Expected results: 101.3 / 96.08 (validation / testing)

```
python3 main_run.py --main-model simple-lstm --epochs 100 --batch-size 80
--dropout 0.15 --nlayers 2 --bptt 70 --nhid 1150 --lr-decay 1 --data
/data/noemien.kocher/datasets/wikitext-2 --shuffle-full-seq
```

Inter-batch TOI

Expected results: 101.7 / 96.89 (validation / testing)

```
python3 main_run.py --main-model simple-lstm --epochs 100 --batch-size 80
--dropout 0.15 --nlayers 2 --bptt 70 --nhid 1150 --lr-decay 1 --data
/data/noemien.kocher/datasets/wikitext-2 --shuffle-row-seq
```

Standard TOI

Expected results: 98.85 / 93.15 (validation / testing)

```
python3 main_run.py --main-model simple-lstm --epochs 100 --batch-size 80
--dropout 0.15 --nlayers 2 --bptt 70 --nhid 1150 --lr-decay 1 --data
/data/noemien.kocher/datasets/wikitext-2
```

Alleviated TOI {2,5,7,10}

Expected results (validation / testing):

- 2: 100.4 / 94.49
- 5: 113.5 / 106.1
- 7: 98.25 / 92.77
- 10: 151.0 / 135.1

```
overlaps=(2 5 7 10)
epochs=100
for k in "${overlaps[@]}"
do
    :
    python3 main_run.py --main-model simple-lstm --epochs
"$(($epochs/$k))" --batch-size 80 --dropout 0.15 --nlayers 2 --bptt 70 --
nhid 1150 --lr-decay 1 --data /data/noemien.kocher/datasets/wikitext-2 --
init-seq "overlapCN_${k}"
    sleep 10
done
```

MOS PTB

Standard TOI:

Expected results: 58.49 / 56.19 (validation / testing)

```
python3 main_run.py --main-model mos-lstm --data data/penn --dropouti 0.4
--dropoutl 0.29 --dropouth 0.225 --seed 28 --batch-size 12 --lr 20.0 --
epochs 1000 --nhid 960 --nhidlast 620 --emsize 280 --n-experts 15
```

Alleviated TOI {2,5,7,10}:

Expected results (validation / testing):

- 7: 57.34 / 55.09

```
python3 main_run.py --main-model mos-lstm --data data/penn --dropouti 0.4
--dropoutl 0.29 --dropouth 0.225 --seed 28 --batch-size 12 --lr 20.0 --
epochs 1000 --nhid 960 --nhidlast 620 --emsize 280 --n-experts 15 --init-
seq overlapCN_7
```

Voice simple LSTM

Extreme TOI:

Expected result: 0.475 / 0.377 (WA / UA)

```
python3 main_run.py --main-model emotions-simple-lstm --cv 5 --data
data/IEMOCAP/all_features_cv --test-batch-size 20 --lr 0.05 --log-interval
20 --lr-decay 1 --step-size 0.1 --epochs 60 --order complete_random
```

Inter-batch TOI:

Expected result: 0.478 / 0.386 (WA / UA)

```
python3 main_run.py --main-model emotions-simple-lstm --cv 5 --data
data/IEMOCAP/all_features_cv --test-batch-size 20 --lr 0.05 --log-interval
20 --lr-decay 1 --step-size 0.1 --epochs 60 --window-size 300 --order
local_order
```

Standard TOI:

Expected result: 0.486 / 0.404 (WA / UA)

```
python3 main_run.py --main-model emotions-simple-lstm --cv 5 --data
data/IEMOCAP/all_features_cv --test-batch-size 20 --lr 0.05 --log-interval
20 --lr-decay 1 --step-size 0.1 --epochs 60 --order standard_order
```

Alleviated TOI 10:

Expected result:

- 15k steps: `0.553` / `0.489` (WA / UA)
- 60 epochs: `0.591` / `0.523` (WA / UA)

```
python3 main_run.py --main-model emotions-simple-lstm --cv 5 --data
data/IEMOCAP/all_features_cv --test-batch-size 20 --lr 0.05 --log-interval
20 --lr-decay 1 --step-size 0.1 --epochs 60 --order total_order
```

Delayed-reset standard TOI {1,2,5,7,10} with PTB

Expected results (validation / testing):

- 1: `61.28` / `58.94`
- 2: `60.76` / `58.55`
- 5: `60.10` / `57.83`
- 7: `60.08` / `57.76`
- 10: `60.05` / `57.78`

```
P=(1 2 5 7 10)
epochs=1000
for k in "${P[@]}"
do
    :
    python3 main_run.py --main-model awd-lstm-repetitions --batch-size 20
--data data/penn --dropouti 0.4 --dropouth 0.25 --seed 141 --epochs 1000 --use-repetitions "${k}"
    sleep 10
done
```

Acknowledgements

Code is heavily borrowed from the following sources:

- simple-lstm (`simple/`): https://github.com/deeplearningathome/pytorch-language-model
- awd-lstm (`awd/`): https://github.com/salesforce/awd-lstm-lm
- mos-lstm: (`mos/`) https://github.com/zihangdai/mos